Funda : ✓

Techniques : → local Nodes
→ Distributed Systems.

→ NFS. (idempotency)

→ Leases
concurrency failure. ✓
* locks can't do the same
* includes time-leases.

NFS : "Distributed" in 1980's.
1990's

↓

client, servers.

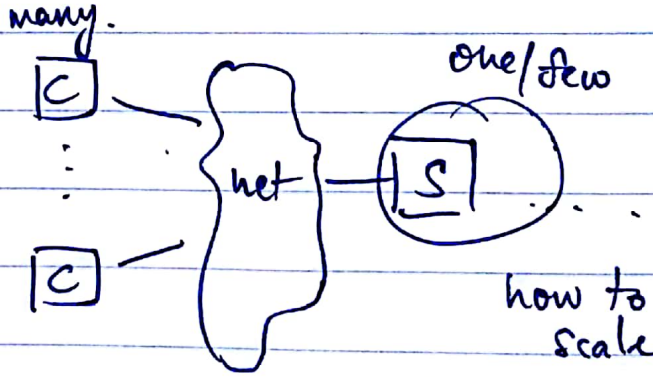


how to scale?

1) scale out (add servers)
→ hard → increases complexity.

happened later (challenging).

2) scale up (make server more powerful).
↑CPU, ↑memory.
easy → fewer interactions

Remzi's opinion.

Main Goal :

goals :-

⇒ transparency.
(performance ?
semantics ? ).

deleting open files and writing to them.

⇒ simplify failures. handling (server).

fault into a performance blip.

[client behaviour. → just try again if anything bad happens]
(idempotency)

$$A = \frac{MTTF}{MTTF + MTTR}$$

↓

reduce to ↑ availability.

[unifies packet loss, server crash] .

based upon :

"statelessness".

why ? open files. handling

local → no problem ; enter system fails

c/s model → partial failure

Protocol:

handle.
(file ~~descriptor~~)

↓

(vol #, inode #, gen #)

- vol # → which file sys?
- inode # → which file?
- gen # → useful why? (across file delete + reuse / inode #).

Ops:

→ read (fh, offset, size).

⇒ data, error code.

→ write (fh, [offset], size, data)

local → os takes care of offset of write.

→ lookup (parent fh, name).

⇒ fh of "name" in parent dir.

→ getattr (fh)

→ stat() syscall about a file.

## from protocal to FS :

not a one to one mapping.

⇒ fs API != DFS protocol.

↓ distributed

### APIs

open ( ) ⟶ lookup

/x/y/z

fh1 . fh2 . . fh3.

read ( ) ⟶ read (s). } one or many.

write ( ) ⟶ write (s).

close ( ) ⟶ . . . (N/A).

only client-side

### How to handle failure?

1) C ⟶ (x) . . . S — just retry

req loss

2) C ⟶ S — retry should be okay.

x ⟵ reply loss (idempotency)

server got smarter. (Server) They cached that they're done it and sends ACK.

3) C ⟶ S (crash)
×

Ans: retry.

Be careful to retry:

→ while (1) {
do:
if success():
break
}

too many retrys causing system to overload.

⇒ [backoff mechanism needed]
required.
↓
wait longer and longer after each retry.

Performance problems:

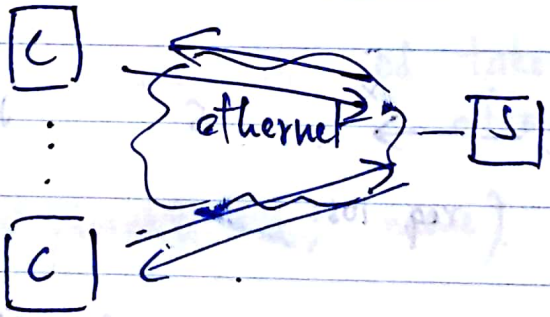


NFS

read caching to prevent server access everything.
[client-side caching]

+ write buffering.

send bulk requests from client

helpful? 1) one chunk of requests.
2) redundant write or actions

Problems:

1)

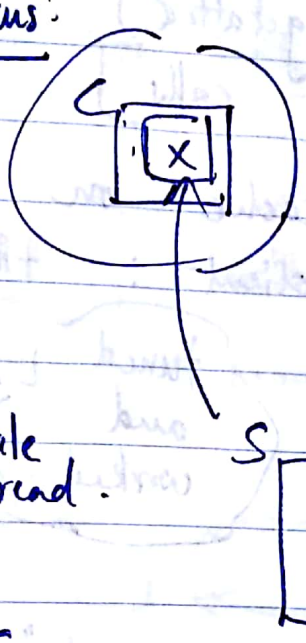


2)

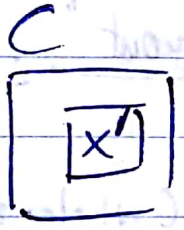


update x on local side

cache consistency

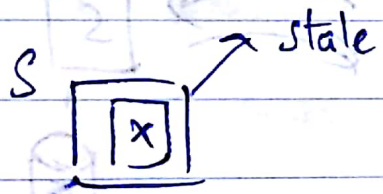


Solutions:

1) C before use of cached: ask if file changed.? get attr() . S . Check before use.

2) flush-on-close:

open
write
:
close → force update on close. to server

last problems :

[too many getattr() calls].

=) attribute cache on client : timeout after

tuned and worked. => [3 seconds].

Leases : "locks w/ timeout" :

AFS example :

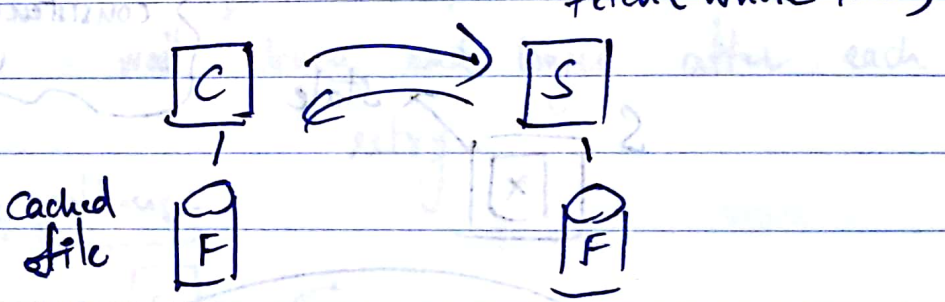


cached file on local disk.

Reads & Write on local.

on close :- if file has changed flush to server

cache on mem or disk ?

↓

better

- persistant (on crash)
- bigger.

AFS v1: cache validation ?

Test Auth : is this file ~~cache~~ valid ?

AFS v2: callback.

promise : server notifies client with changes to file.

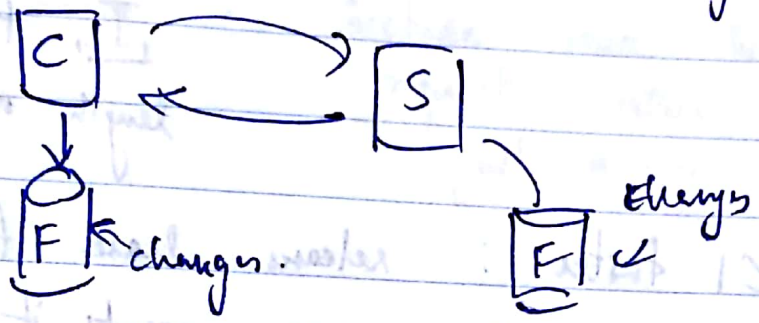


with a cache: not stateless anymore.

Server crash handling is different.

=) complex.

client crash =) how to notify.

Solution:

leases rights with timeout.

⇓

key to handling failure.

common:

acquire → use it

(maybe renew periodically)

release (or just let timeout)

=) when safely grant lease to another client ?

clocks.

c1

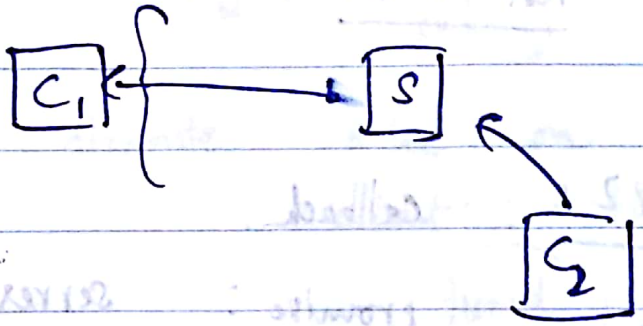


wants clock C2.

1) server waits, wait how long?

T + Δ clock skew that we tolerate.

length of lease.

Clock { C1 faster : releases lease faster
C1 slow : server grants it to C2 and C1 thinks it has lease.

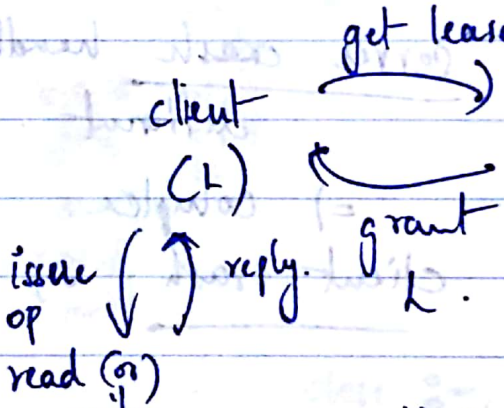


lease: grants exclusive access to block/range

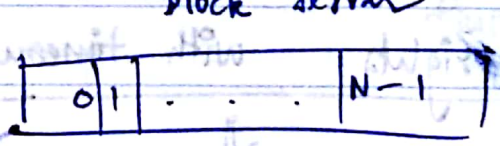


leases to make sure ops ~~are~~ to block server is serial.

Problem: Race condition

{ client requests even after lease expires on server. and S grants lease to another client.

not serial anymore.

how to handle?:

2) fence : include lease info in request. rather than just relying on timing.